

Vishwa Swaroop Padarti

✉ visuvishwa7227@gmail.com | 📞 508-630-6815 | 🏠 Phoenix, AZ
🌐 www.linkedin.com/in/vishwa-p-109288191

PROFESSIONAL SUMMARY

- ▷ **Senior Data Engineer** with **9+ years** of experience building cloud data platforms across finance, retail, telecom, and EdTech. Specialized in legacy-to-cloud migrations, lakehouse modernization, batch and streaming pipelines, data modeling, observability, and cost optimization across **AWS, Azure, Databricks, Snowflake, Redshift, BigQuery, and dbt**.
- ▷ **Certified Azure and AWS data specialist** experienced in delivering governed, reliable, and analytics-ready platforms for business-critical reporting and enterprise data modernization, with current focus on AI-assisted data operations using governed cloud data tools.

SKILLS

Languages: *Python, Scala, SQL, PL/SQL, Shell Script*
Cloud Platforms: *AWS (EMR, S3, Glue, Lambda, Redshift, Athena, Step Functions, SQS, EventBridge), Azure (ADF, ADLS Gen2, Event Hubs, Key Vault), GCP (BigQuery, GCS)*
Modern Data Stack: *Databricks, Snowflake, dbt Cloud/dbt Mesh, Delta Lake, Iceberg, Medallion Architecture, Data Modeling, Fivetran, CDC, Reverse ETL, Airflow*
Big Data: *Apache Spark, Kafka, Hive, NiFi, Airflow*
Agentic Data Ops: *Bedrock AgentCore, Strands Agents, Athena Tools, CloudWatch Metrics, Glue Job Monitoring, LangChain/RAG*
DevOps: *Docker, Kubernetes, CI/CD, GitHub, Bitbucket*
Visualization & Governance: *Power BI, Tableau, Collibra, Great Expectations*

WORK EXPERIENCE

Lexia Learning Aug 2025 — Present

Senior Data Engineer (Consultant, Remote)

- Improved post-acquisition reporting trust by migrating **Lexia** learning-product data from **MySQL to Redshift** and operationalizing governed **AWS Glue + dbt Mesh** pipelines for curated analytics across Core5, LETRS, and Aspire.
- Reduced operational overhead for student-activity ingestion by building a serverless event platform using **SNS/SQS FIFO, EventBridge Pipes, Kinesis Firehose, S3, Glue, and Athena Iceberg**.
- Accelerated pipeline triage and data debugging by implementing a **Bedrock AgentCore + Strands Agents** pipeline intelligence layer with guarded **CloudWatch, Glue, and Athena** tools for natural-language health checks and read-only analytics.
- **Key Technologies:** *AWS (S3, Redshift, Glue, SNS/SQS, EventBridge, Firehose, Athena, CloudWatch, Bedrock AgentCore), MySQL, dbt, Strands Agents, Workato, Tableau, CI/CD Pipelines.*

CAVA Jul 2024 — Jul 2025

Senior Data Engineer (Consultant)

- Reduced pipeline processing time by **30%** by consolidating **Brink POS, Crunchtime, NetSuite**, and operational feeds into **Databricks Delta** Bronze/Silver lakehouse layers for store operations analytics.
- Reduced cloud compute cost by **20%** and improved query performance by **25%** by optimizing Redshift-to-Databricks migration patterns, Delta layouts, and store analytics workloads across **300+** locations.
- Improved restaurant operations data reliability by building metadata-driven **Oracle Crunchtime** ingestion with JDBC, S3 Parquet, schema validation, audit logging, primary-key dedupe, and soft-delete reconciliation.
- Enabled replay-safe POS analytics by implementing **Brink SOAP/XML GetOrders** ingestion, raw XML preservation, nested order flattening, and timestamp-guarded Delta merges.
- **Key Technologies:** *AWS (S3, Redshift, Managed Airflow, KMS, Glue), Databricks, Delta Lake, Databricks Asset Bundles, Medallion Architecture, Data Modeling, Docker, SQL, Python, DataDog, Boomi, Tableau, CI/CD, Brink POS, Crunchtime, NetSuite.*

Senior Data Engineer (Consultant, Remote)

- Migrated **250TB** of historical Oracle/ERP data into Azure and Snowflake-backed analytics layers, improving data accessibility by **20%** across **8** business-critical pipelines.
- Maintained **99.9%** data integrity by implementing **Great Expectations** quality checks and GitHub-based automated error detection across **ADF, ADLS Gen2, Snowflake, and dbt** workflows.
- Reduced operational issue detection time from **4 hours to 30 minutes** by building **Kafka/Event Hubs** pipelines and **PySpark** monitoring frameworks across **3** data products.
- **Key Technologies:** *Azure (ADF, ADLS Gen2, Event Hubs, Key Vault), Snowflake, Databricks, dbt Cloud, Great Expectations, Data Modeling, Docker, Kubernetes, SQL, Python, Power BI, Apache (Kafka, Spark), CI/CD Pipelines.*

Freddie Mac

Dec 2021 — Dec 2022

AWS Data Engineer (Consultant, Remote)

- Reduced mortgage data processing runtime by **25%** by refactoring **PySpark** transformations and orchestration for ULAD Mongo raw/raw+ pipelines on **AWS EMR**.
- Processed **3.5M** mortgage applications monthly by developing the **BYOL/EDP** framework with EMR Spark, Step Functions, DynamoDB control metadata, Glue, S3, and Snowflake.
- Improved downstream analytics reliability by shredding denormalized loan-evaluation documents into **8 enhanced staging datasets** and **18 trusted/optimized tables** for regulated reporting workloads.
- **Key Technologies:** *AWS (DynamoDB, S3, Glue, Step Functions, EMR), MongoDB, JSON, PySpark, Python, SQL, Snowflake, Data Governance (Collibra), CI/CD Pipelines*

American Express

Nov 2018 — Nov 2021

Big Data Engineer II (Consultant)

- Improved risk assessment accuracy by **3%** by building ETL workflows in **Magellan, Lucy, Hive, Oozie, and Tez** for fraud detection, eligibility, and proactive credit-line decisioning.
- Evaluated BigQuery modernization for **50TB+** of transaction and risk data by migrating legacy Lucy/Hive business logic to **GCP BigQuery** and comparing query performance/cost against existing Hive workflows.
- Strengthened audit-ready operations for regulated financial datasets through data quality controls, secure access patterns, and governance-focused pipeline support.
- **Key Technologies:** *ETL (Magellan, Lucy), Apache (Hive, Oozie, Tez, ORC), GCP (BigQuery), REST API, Cornerstone, CI/CD Pipelines.*

Global Atlantic Financial Group

Oct 2017 — Oct 2018

Data Engineer (Consultant)

- Reduced financial ingestion runtime by **72%** (from **17 hours to 5 hours**) by tuning Sqoop jobs and standardizing **NiFi/Spark** ingestion into secured HDFS encryption zones.
- Delivered SLA-backed warehouse ingestion patterns across SFTP and database sources using **Sqoop, NiFi, Spark, Hive, Ranger, Kerberos, and Oozie**.
- **Key Technologies:** *AWS (EC2, EBS, S3, Redshift, RDS, VPC, CloudWatch), HDP (NiFi, Hive, Sqoop, Spark, Ranger, Oozie), File Format (Text, Avro, ORC), SQL (Oracle, SybaseIQ), SFTP Server, Scala, Python3, Boto, Data Security (Ranger, Kerberos), Data Visualization (Tableau), CI/CD Pipelines*

EDUCATIONMaster's in Computer Science, Bradley University (**Peoria, IL**)

Aug 2014 - May 2016

CERTIFICATIONS

- Microsoft Certified: Azure Data Engineer Associate
- AWS Certified Data Analytics – Specialty
- dbt Fundamentals – dbt Labs
- Databricks Academy (Gen AI Fundamentals)